v3.3.3

OpenOva

The Converged Cloud-Native Platform

Pure open-source. AI-native. Zero lock-in.
Infrastructure, data, AI, and communication — deployed in hours, not years.

Pure Open Source   AI-Native   CNCF Mature   Community = Enterprise

openova.io · hello@openova.io · 2026

OpenOva

Why Now

The $6.15 Trillion Question

Sources: Gartner Feb 2026 · Deloitte · MIT Sloan · CNBC · Zylo · Flexera · IBM IBV

Where $6.15 Trillion Goes

$1.9T IT Services — 31%

$1.4T Software — 23%

$1.4T Telecom — 23%

$836B Devices — 14%

$653B Data Ctrs — 11%

Global IT spending reaches **$6.15 trillion** in 2026 according to Gartner's latest forecast. The number sounds like progress, but beneath it lies a staggering inefficiency that threatens to derail the AI revolution before it begins. Enterprises are spending more than ever, yet **75 cents of every dollar** goes to maintenance: patching servers, renewing licenses, managing vendors, servicing technical debt accumulated over decades.

The average large enterprise now runs **660 SaaS applications** across departments, with organizations above 10,000 employees averaging $284 million in annual SaaS spend alone. Most of these tools were adopted bottom-up — individual teams purchasing subscriptions with zero centralized oversight. The result is devastating: **53% of all software licenses** sit completely unused, a silent tax bleeding budgets dry while delivering nothing (Zylo 2025 SaaS Management Index).

The 75-Cent Problem

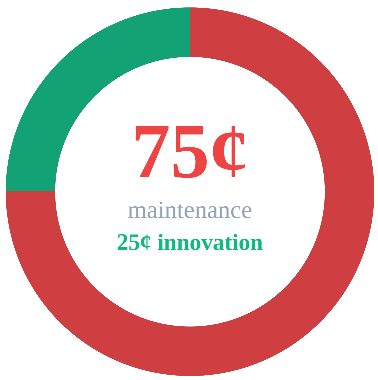

For every dollar an enterprise spends on IT, **75 cents** goes to keeping the lights on: patching operating systems, managing vendor relationships, maintaining legacy integrations, renewing licenses for software nobody uses. Only **25 cents** funds anything genuinely new. Your highest-paid engineers spend their days as plumbers and janitors rather than architects and builders.

This isn't just a budget problem — it's an **AI readiness crisis**. **72% of CIOs** now identify fragmented infrastructure as the number-one barrier to AI ROI (Gartner 2025). The models work — GPT-4, Claude, Gemini all perform brilliantly in isolation. But when enterprises try to deploy them against real workflows, they slam into siloed data lakes, incompatible APIs, missing observability, and zero governance frameworks.

The proof is damning. **95% of GenAI pilot programs** fail to deliver measurable business value (MIT Sloan Management Review, 2025). Not because the algorithms are flawed, but because the infrastructure substrate beneath them is fragmented beyond repair. AI needs unified data, consistent APIs, and platform-wide observability. Instead, it gets 660 disconnected SaaS tools, 3 competing cloud providers, and 14 different authentication systems.

The enterprise software industry has spent two decades optimizing for **vendor lock-in** rather than interoperability. Every product decision prioritized switching costs over developer experience. The chickens are now coming home to roost — and the market knows it.

The Market is Speaking — Q1 2026

-$1T Market Cap Lost

-$285B Single-Day Wipeout

+1500% VMware License Hike

Investors are punishing fragmented, legacy-dependent software vendors at historic levels. In Q1 2026 alone, **$1 trillion** in software market capitalization was destroyed as Wall Street concluded that the old guard cannot deliver AI-ready infrastructure (CNBC). A single trading day in January saw **$285 billion** wiped from AI sector stocks after investors realized hype without substance gets punished swiftly and brutally (WinBuzzer).

The vendor shakeout is accelerating across every segment. Project management SaaS stocks are down **58% year-to-date** as enterprises question why they pay premium prices for tools that don't integrate. IT outsourcing firms have lost **42%** of their market value. Proprietary CRM vendors are down **38%**. Virtualization incumbents have shed **27%** — investors see the writing on the wall as open-source alternatives mature.

Meanwhile, proprietary vendors are making the problem worse. VMware's post-acquisition license restructuring hit customers with increases up to **+1,500%**, forcing pay-or-migrate ultimatums on captive customers with no escape route (CIO/Avasant). This is extraction, not innovation. Smart money is betting accordingly: hedge funds made **$24 billion** in profits shorting legacy enterprise vendors in just five weeks (CNBC).

The message from capital markets is unambiguous: **fragmented infrastructure is a liability, not an asset**. Enterprises that continue bolting together dozens of incompatible tools will face rising costs, Cloud waste compounds the problem. **27% of the $450 billion+** global cloud infrastructure spend in 2025 is pure waste: idle instances running 24/7 with no workloads, over-provisioned virtual machines sized for peak loads that never arrive, and forgotten resources from projects abandoned months ago. Per organization, that translates to roughly **$17 million per year** burned with zero return (Flexera 2025 State of the Cloud). When organizations try to fix this sprawl, they hit another wall: **94% of core modernization projects** exceed their timelines and budgets (IBM Institute for Business Value). Complexity kills transformation before it starts.

falling productivity, and AI programs that never move beyond pilot stage. The question isn't whether to transform — it's how fast you can converge before your competitors do.

Option A

OpenOva

The Industry Response

Five Megatrends. One Convergence.

Five megatrends — each backed by billions in investment and analyst consensus — all point to one architecture.

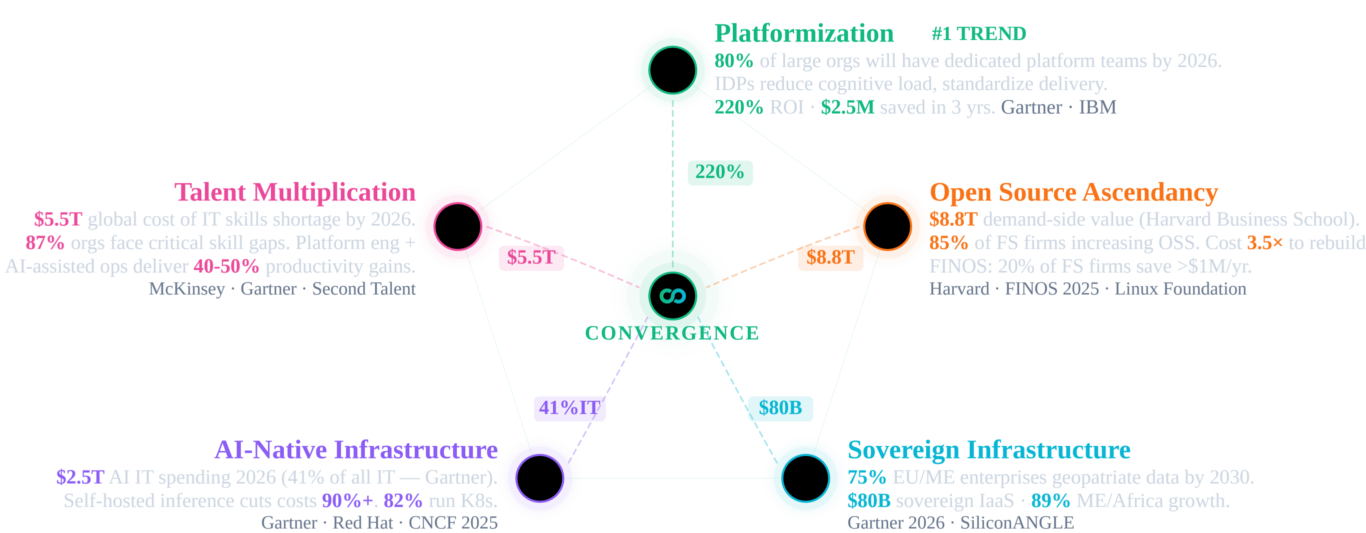

OpenOva sits at the intersection of all five. A full-stack, open-source, AI-native platform — with enterprise support. Born cloud-native. Zero lock-in.

Platform Architecture

Building Blocks — Concept View

Ten building blocks, one converged ecosystem. Each solves a distinct infrastructure challenge — independently adoptable, collectively transformative.

One platform, ten domains. Each building block solves a distinct infrastructure challenge — security, networking, observability, data, AI, communication, storage, scaling, GitOps, and lifecycle. Together, they replace the 660-app sprawl typical of large enterprises (10K+ employees — Zylo 2025) with a single converged ecosystem.

Fully pluggable — keep what works. OpenOva doesn't demand rip-and-replace. Already running a database you love? Keep it. Have a working CI/CD pipeline? Plug it in. Each building block is independently adoptable. Start with what hurts most, expand when ready. Your existing investments are preserved.

Containers first, VMs as a last resort. 95% of new digital workloads deploy on cloud-native platforms (Gartner 2025). OpenOva is container-native from the ground up. For the rare legacy workload that can't be containerized, Cloud CRDs (Crossplane) provision and manage VMs through the same GitOps workflow — no separate toolchain needed.

Born cloud-native, AI-native. Every component runs as a Kubernetes operator. GitOps-driven. Observable from day one. AI-ready by default. This is the platformization megatrend — purpose-built, not bolted on.

Convergence without the integration tax. What starts as à la carte adoption converges into a unified ecosystem — shared identity, shared observability, shared security — zero glue code, zero custom integrations.

AI needs clean infrastructure. 72% of CIOs can't get ROI from AI — not because the models are bad, but because fragmented data, siloed observability, and disconnected security prevent AI from delivering. OpenOva unifies the substrate so AI can actually work.

Security
TLS Automation Secrets Sync Secrets Vault Vuln Scan Runtime Security Supply Chain Trust
SBOM & CVE Policy Engine IAM
AI Hub
Model ServingServerlessLLM Inference Vector DBGraph DBEmbeddings Chat UIAI GuardrailsLLM Observability LLM Gateway
Data & Integration
PostgreSQLMongoDB WireRedis Cache Event StreamingAnalytics DBStream Processing
Workflow EngineCDCData Lakehouse BI & DashboardsOpen Banking
Communication
Email ServerVideo & AudioWebRTC GWChat ProtocolPush Notify
GitOps & IaC
GitOps EngineGit ServerIaCCloud CRDs
Networking & Mesh
CNI & MeshL7 ProxyWAFDNS Sync GSLBReverse TunnelMesh VPNIPsec Gateway
Scaling & Resilience
Vertical ScalingEvent ScalingConfig ReloadHA Orchestration
Storage & Registry
Object StorageBackup & RestoreContainer Registry
AIOps
AIOps Brain Dashboards Telemetry Agent Log Aggregation Metrics Store Distributed Tracing
Instrumentation Search & Analytics Chaos Engineering Usage Metering
Bootstrap & Lifecycle
Initial Setup Day-2 Operations
Kubernetes (K3s)
Huawei Cloud
AWS
Oracle OCI
Hetzner

Ten building blocks. One converged ecosystem. Start with what hurts most — expand when ready.

Open Source Foundation

Building Blocks — Tech View

Every component is real upstream open source — CNCF-graduated, battle-tested, community-backed. No proprietary forks. No rebranded wrappers.

$8.8 trillion in demand-side value. A 2024 Harvard Business School study calculated that open-source software creates $8.8T in demand-side economic value globally — meaning if organizations had to rebuild or license equivalent proprietary software, the cost would be 3.5× higher. 85% of financial services firms are increasing OSS adoption (FINOS 2025). 20% already save >$1M/year.

CNCF-graduated, enterprise-proven. Every OpenOva component — Cilium, Falco, Flux, Grafana, Strimzi, KServe, Keycloak — is selected for production maturity and backed by 15.6M cloud-native developers. No rebranded forks. No proprietary wrappers. Real upstream projects.

Enterprise support, zero lock-in. Open source doesn't mean unsupported. OpenOva provides 24/7 enterprise support, SLAs, and expert guidance — while you keep full ownership of every line of code. Switch providers without switching technology. Community edition = Enterprise edition.

Open source is more secure, not less. The largest breaches of the decade all came from proprietary, closed-source software: a supply-chain attack on a network monitoring platform ($100B+ impact), an endpoint agent update that crashed 8.5M devices worldwide, a file-transfer zero-day exploited across thousands of organizations ($10B+ damages), a password manager breach exposing encrypted vaults, and a collaboration platform exploit affecting government agencies. Meanwhile, open-source code is audited by millions, with CVEs patched in hours. Intentional backdoors can't hide in code the world can read.

GUARDIAN

cert-manager ESO OpenBao Trivy Falco Sigstore Syft+Grype Kyverno Keycloak

CORTEX

KServeKnativevLLM MilvusNeo4jBGE LibreChatNeMoLangFuse Axon

FABRIC

CNPGFerretDBValkey StrimziClickHouseFlink TemporalDebeziumIceberg SupersetFingate

RELAY

StalwartLiveKitSTUNnerMatrixNtfy

PILOT

FluxGiteaOpenTofuCrossplane

SPINE

CiliumEnvoyCorazaExternalDNS k8gbfrpcNetBirdstrongSwan

SURGE

VPAKEDAReloaderContinuum

SILO

MinIOVeleroHarbor

INSIGHTS

Specter Grafana Alloy Loki Mimir Tempo OpenTelemetry OpenSearch Litmus OpenMeter

CATALYST

Bootstrap Wizard Lifecycle Manager

Kubernetes (K3s)

Huawei Cloud

AWS

Oracle OCI

Hetzner

Community IS enterprise. Same code, same features, same security — with 24/7 commercial support.

OpenOva

Portfolio Spotlight

CATALYST: Bootstrap & Lifecycle

From initial platform setup to Day-2 operations — guided, automated, and continuously managed. Runs on K3s across any cloud.

Platform engineering is the #1 strategic technology trend (Gartner 2024), yet 94% of core modernization projects exceed timelines and budgets (IBM IBV). The bottleneck isn't technology — it's the operational complexity of assembling, configuring, and maintaining dozens of cloud-native components across environments.

CATALYST eliminates this complexity. A guided bootstrap wizard provisions a production-ready platform in hours, not months. Day-2 lifecycle management — automated upgrades, health monitoring, drift repair — keeps every component current and healthy without manual intervention. Deploy on any cloud or bare metal, from edge to enterprise.

Bootstrap Wizard

Guided platform provisioning with interactive configuration and automated deployment

Lifecycle Manager
Continuous Day-2 operations with automated upgrades, health checks, and drift repair

K3s
Lightweight Kubernetes distribution optimized for edge and production workloads

Cloud Providers
Multi-cloud support: Huawei Cloud, AWS, Oracle OCI, Hetzner, and bare metal

Production-grade Kubernetes deployed in hours, not months. Automated Day-2 operations from day one.

Portfolio Spotlight

PILOT: GitOps & IaC

Git-driven delivery and infrastructure as code. Every change is auditable, reversible, and automated.

The shift to GitOps is accelerating — 76% of organizations now use or evaluate GitOps practices (CNCF 2025). Yet many rely on proprietary Git hosting and IaC tools with restrictive licensing changes. License compliance risk is real: the BSL/SSPL trend forces enterprises to re-evaluate their entire delivery toolchain.

PILOT delivers a fully integrated, license-safe GitOps stack. Every infrastructure change is Git-versioned, auditable, and reversible. Crossplane Cloud CRDs extend GitOps beyond Kubernetes — managing cloud resources, VMs, and databases through the same declarative workflow. Zero vendor lock-in on your delivery pipeline.

Flux
Continuous GitOps delivery engine that reconciles cluster state with Git as the single source of truth

Gitea
Lightweight internal Git hosting with built-in CI/CD pipelines and issue tracking

OpenTofu
MPL 2.0 licensed infrastructure-as-code for declarative cloud provisioning and management

Crossplane
Kubernetes-native cloud resource management using custom resource definitions and reconciliation

Every change auditable, reversible, and Git-driven. Infrastructure as code is not optional — it's the foundation.

Portfolio Spotlight

SPINE: Networking & Mesh

eBPF-powered service mesh, L7 proxy, WAF, DNS automation, global load balancing, and secure connectivity.

Enterprise networking is fragmenting across sidecar proxies, proprietary load balancers, managed WAFs, and per-seat VPN licenses. Each adds latency, cost, and operational burden. Traditional service meshes inject sidecar containers into every pod — adding 15-30% memory overhead and significant complexity.

SPINE leverages eBPF to eliminate sidecars entirely. Cilium provides CNI, service mesh, mTLS, and L7 observability in a single kernel-level data plane. Add WAF protection, automated DNS, global load balancing, and secure connectivity across hybrid and edge environments — all managed declaratively through Kubernetes.

Cilium
eBPF-powered CNI with built-in service mesh, mTLS, and L7 visibility

Envoy
High-performance L7 proxy for traffic management, load balancing, and observability

Coraza
Open-source web application firewall implementing OWASP Core Rule Set protection

ExternalDNS
Automated DNS record synchronization between Kubernetes resources and DNS providers

k8gb
Global server load balancing with authoritative DNS for multi-cluster traffic management

frpc
Reverse tunnel proxy for exposing services behind NAT in hybrid and edge deployments

NetBird
Zero-config WireGuard mesh VPN for secure peer-to-peer connectivity across clusters

strongSwan
IPsec VPN gateway for site-to-site encryption and legacy network integration

eBPF-native networking at kernel speed. No sidecars, no overhead — from CNI to service mesh to VPN in one stack.

Portfolio Spotlight

GUARDIAN: Security & Identity

End-to-end security: TLS, secrets, scanning, runtime protection, supply chain integrity, policy enforcement, and identity management.

The largest breaches of the decade came from proprietary, closed-source software — a supply-chain attack on a network monitoring platform ($100B+ impact), an endpoint agent update that crashed 8.5M devices, a file-transfer zero-day ($10B+ damages). Meanwhile, open-source code is audited by millions. Backdoors can't hide in code the world can read.

GUARDIAN wraps your platform in defense-in-depth: automated TLS, zero-trust secrets, eBPF runtime detection, supply chain signing, SBOM tracking, policy-as-code, and enterprise IAM. Every layer is open source, auditable, and integrated — from build pipeline to production runtime to identity.

cert-manager
Automates TLS certificate lifecycle with Let's Encrypt and private CA integration

ESO
Syncs secrets from external vaults into Kubernetes with automatic rotation support

OpenBao
MPL 2.0 secrets vault with per-cluster isolation and dynamic credential generation

Trivy
Comprehensive vulnerability scanner for container images, filesystems, and IaC templates

Falco
eBPF-based runtime threat detection for containers and Kubernetes workloads

Sigstore
Keyless container signing and verification for supply chain integrity assurance

Syft+Grype
Automated SBOM generation paired with vulnerability matching and CVE tracking

Kyverno
Policy engine for validating, mutating, and generating Kubernetes resources at admission

Keycloak
Enterprise identity and access management with OIDC, SAML, and FAPI support

Defense-in-depth from build to runtime. Zero-trust by default — mTLS, RBAC, policy-as-code, fully auditable.

OpenOva

Portfolio Spotlight

INSIGHTS: AIOps & Observability

Full observability stack with AI-native operations. Specter's 6 autonomous agents replace legacy SOAR/SIEM — DevSecOps, DevOps, SRE, FinOps, Compliance, and AI Ops.

Enterprises spend $300-500K/yr on proprietary observability — per-host pricing that scales linearly with infrastructure. As cloud-native environments grow to thousands of containers, cost becomes prohibitive. Meanwhile, alert fatigue and siloed tools mean MTTR stays high despite massive spending.

INSIGHTS unifies logs, metrics, traces, and SIEM into a single stack with zero per-host fees. Specter's 6 autonomous AI agents — DevSecOps, SRE, FinOps, Compliance, DevOps, and AI Ops — replace the traditional SOC/NOC model. They detect, diagnose, and remediate issues before your team even wakes up. Chaos engineering and usage metering round out a complete operational platform.

Specter
Six autonomous AI agents for DevSecOps, SRE, FinOps, and continuous compliance

Grafana
Unified dashboards for metrics, logs, traces, and alerts in a single pane of glass

Alloy
Next-generation OpenTelemetry collector and telemetry pipeline agent by Grafana Labs

Loki
Horizontally scalable log aggregation system optimized for Kubernetes metadata queries

Mimir
Long-term metrics storage with global query capabilities and multi-tenancy support

Tempo
Distributed tracing backend with native OpenTelemetry support and deep Grafana integration

OpenTelemetry
Vendor-neutral instrumentation framework for automatic trace and metric collection

OpenSearch
Search and analytics engine powering hot-tier SIEM and security event correlation

Litmus
Chaos engineering platform for proactive resilience testing of Kubernetes workloads

OpenMeter
Real-time usage metering and billing engine for consumption-based pricing models

Six AI agents. Zero per-host fees. Detect, diagnose, and fix autonomously — unified logs, metrics, traces, and SIEM.

Portfolio Spotlight

SURGE: Scaling & Resilience

Intelligent autoscaling, configuration management, and high-availability orchestration.

Over-provisioning wastes 27% of cloud spend ($44.5B globally), while under-provisioning causes outages. Traditional scaling is reactive — you pay for peak capacity 24/7 or risk downtime. Manual capacity planning fails in dynamic, event-driven architectures.

SURGE combines vertical right-sizing (VPA), event-driven horizontal scaling (KEDA), and continuous availability orchestration into an intelligent autoscaling engine. Workloads scale from 10 to 10,000 pods based on actual demand — queue depth, HTTP traffic, custom metrics. Zero-downtime deployments with automatic failover and configuration-change detection.

VPA
Automatically right-sizes pod CPU and memory requests based on actual utilization patterns

KEDA
Event-driven horizontal autoscaling triggered by queue depth, HTTP traffic, or custom metrics

Reloader
Watches ConfigMaps and Secrets for changes and triggers rolling restarts automatically

Continuum
Continuous availability orchestration ensuring zero-downtime deployments and failovers

Scale from 10 to 10,000 pods on demand. Smart autoscaling eliminates the $44.5B global cloud waste problem.

Portfolio Spotlight

SILO: Storage & Registry

S3-compatible object storage, backup & disaster recovery, and container artifact registry.

Data sovereignty demands self-hosted storage. Cloud egress fees ($0.08-0.12/GB) and proprietary registry lock-in make multi-cloud and hybrid deployments expensive and inflexible. Backup strategies that depend on a single cloud provider fail the sovereignty test.

SILO provides S3-compatible object storage, Kubernetes-native backup with cross-region replication, and a full container/Helm registry — all self-hosted. Run it on-prem, in your sovereign cloud, or alongside public clouds. Your data stays where regulations and business logic dictate. No egress surprises.

MinIO
S3-compatible high-performance object storage for cloud-native data and backup targets

Velero
Kubernetes-native backup, restore, and disaster recovery with cross-region replication

Harbor
Enterprise container and Helm chart registry with vulnerability scanning and RBAC

Your data stays where regulations demand. S3-compatible object storage, cross-region DR — zero egress surprises.

Portfolio Spotlight

CORTEX: AI Hub

Self-hosted AI infrastructure. Full control over models, prompts, and costs. No data leaves your environment.

$2.5T will be spent on AI in 2026, with 90% going to inference (Gartner). Yet 95% of GenAI pilots fail ROI (MIT) — not because of the models, but because enterprises lack the infrastructure to run them. API-based AI services leak data, impose per-token fees, and create hard vendor lock-in on the most strategic technology of the decade.

CORTEX gives you full control. Self-hosted multi-model inference cuts token costs by 90%+. RAG pipelines with vector and graph databases. AI safety guardrails. LLM observability. A ChatGPT-like interface for business users. All running on your infrastructure, with your data, under your governance. Switch models freely — no API lock-in.

KServe
Kubernetes-native model serving with autoscaling, canary rollouts, and multi-framework support

Knative
Serverless platform for deploying event-driven and scale-to-zero inference workloads

vLLM
High-throughput inference engine with PagedAttention for optimal GPU utilization

Milvus
Distributed vector database for similarity search in RAG and recommendation pipelines

Neo4j
Graph database for knowledge graphs, entity relationships, and AI-powered reasoning

BGE
State-of-the-art embedding and reranking models for retrieval-augmented generation

LibreChat
ChatGPT-like interface supporting multiple LLM backends for enterprise AI chat

NeMo
AI safety guardrails framework for input/output filtering and responsible AI deployment

LangFuse
LLM observability platform with prompt tracing, cost analytics, and evaluation tools

Axon
Managed LLM gateway providing subscription-based inference access and rate limiting

Self-hosted AI inference cuts costs 90%+. Run any model on your infrastructure — no data leaves your environment.

Portfolio Spotlight

FABRIC: Data & Integration

Event-driven data platform: PostgreSQL, streaming, workflows, CDC, data lakehouse, and BI analytics.

Data infrastructure is the most expensive line item in enterprise IT. Proprietary database licensing alone costs $100-300K/yr per organization, while managed event streaming adds another $80-120K. Per-core pricing punishes scale. License audits punish growth. And each vendor creates its own data silo.

FABRIC replaces fragmented data infrastructure with a unified, open-source data platform. Cloud-native PostgreSQL with HA and automated backups. Event streaming for real-time architectures. Workflow orchestration for distributed transactions. A complete data lakehouse for analytics. Self-service BI. All integrated, all open source, all under your control.

CNPG
Cloud-native PostgreSQL operator with HA, automated backups, and declarative management

FerretDB
MongoDB wire protocol compatibility layer running on PostgreSQL for document workloads

Valkey
Redis-compatible in-memory cache and data store with full cluster mode support

Strimzi
Apache Kafka operator for production event streaming with automated topic management

ClickHouse
Columnar OLAP database for real-time analytics on billions of rows with SQL interface

Flink
Unified stream and batch processing engine for complex event processing pipelines

Temporal
Durable workflow orchestration for sagas, long-running processes, and distributed transactions

Debezium
Change data capture platform streaming database changes to Kafka in real time

Iceberg
Open table format enabling data lakehouse architecture with time-travel queries

Superset
Self-service BI platform with interactive dashboards, SQL Lab, and data exploration

Fingate
Open banking platform with FAPI authorization, OBIE APIs, and usage metering

PostgreSQL to real-time CDC to BI dashboards. Zero per-core licensing — one unified open-source data platform.

Portfolio Spotlight

FINGATE: Open Banking

End-to-end open banking solution compliant with FAPI and OBIE standards. Built entirely on open-source components from the OpenOva stack.

Open banking mandates (PSD2, OBIE, Saudi Arabia SAMA) require FAPI-certified APIs — yet proprietary platforms charge per-API-call fees that scale exponentially with transaction volume. Banks face a choice: pay ever-growing vendor fees or build from scratch. Neither option is sustainable for the 89% growth in ME/Africa digital banking.

FINGATE delivers a complete open banking solution built entirely on OpenOva's open-source stack. FAPI-certified authorization, OBIE-compliant APIs, real-time metering, payment orchestration with saga patterns, and financial-grade API protection. Zero per-API-call fees. Full regulatory compliance. Built for the next decade of digital finance.

Keycloak (FAPI)
Financial-grade API authorization server with Strong Customer Authentication for PSD2

OBIE APIs
Open Banking Implementation Entity standards for accounts, payments, and consent management

Temporal
Payment orchestration engine for PSD2 consent flows and settlement workflows

OpenMeter
API consumption metering with real-time billing and partner settlement reporting

Envoy + Coraza
API gateway with WAF protection, rate limiting, and financial-grade traffic management

CNPG
PostgreSQL for transactional ledger data with point-in-time recovery and replication

Strimzi + Debezium
Kafka event streaming with CDC for real-time balance synchronization and audit trails

FAPI-certified. PSD2 and SAMA compliant. 100% open source, zero per-API-call fees — built for the ME/Africa boom.

How We Work

Engagement Model

Two service tracks — platform and consultancy — take you from discovery to autonomous operations.

Platform Services

Subscription-based, self-service + managed

Community Edition

Full platform, self-managed. Community = Enterprise — same code, same features.

Enterprise Subscription Agreement (ESA)

24/7 enterprise support, SLAs, Expert Network access (80hr/quarter), continuous updates, quarterly business reviews.

Fully Managed SOC/NOC

Specter AI agents + human oversight. 24/7 autonomous operations — DevSecOps, SRE, FinOps, Compliance.

Pay-As-You-Go (Axon)

Managed LLM inference — consumption-based pricing, no GPU procurement, instant access.

Consultancy Services

Engagement-based, expert-led

Discovery & Strategy

Full landscape audit — we map every vendor contract, license obligation, and integration dependency. TCO/ROI financial modeling. Target architecture design that preserves what works. Board-ready strategy with clear milestones and risk quantification.

Statement of Work (SoW)

Fixed-scope delivery with defined milestones: platform deployment, workload containerization and migration (Exodus methodology), compliance alignment, hands-on team enablement, and knowledge transfer.

Time & Materials (T&M)

Embedded experts that scale with your transformation. Architecture reviews, performance optimization, security hardening, and specialist deep-dives — ramping up or down as phases progress.

CIO & DT Advisory

C-suite strategic partnership: technology roadmap aligned to business KPIs, organizational readiness assessment, change management, transition planning, and quarterly board reporting.

Global Reach

The Guild — Expert Network

Global presence with deep expertise in banking, regulated industries, and enterprise Kubernetes.

Muscat — Headquarters

Corporate headquarters, strategic operations, Middle East market

Dubai — Regional Presence

Financial services, enterprise partnerships, Gulf market

Amsterdam — European Hub

European market, regulatory compliance, enterprise sales

Istanbul — Engineering & Operations

Core development, platform engineering, DevOps, security

Ankara — Engineering Hub

Platform R&D, AI/ML engineering, product development

Bengaluru — Asia-Pacific Hub

Engineering scale, AI/ML talent, APAC market

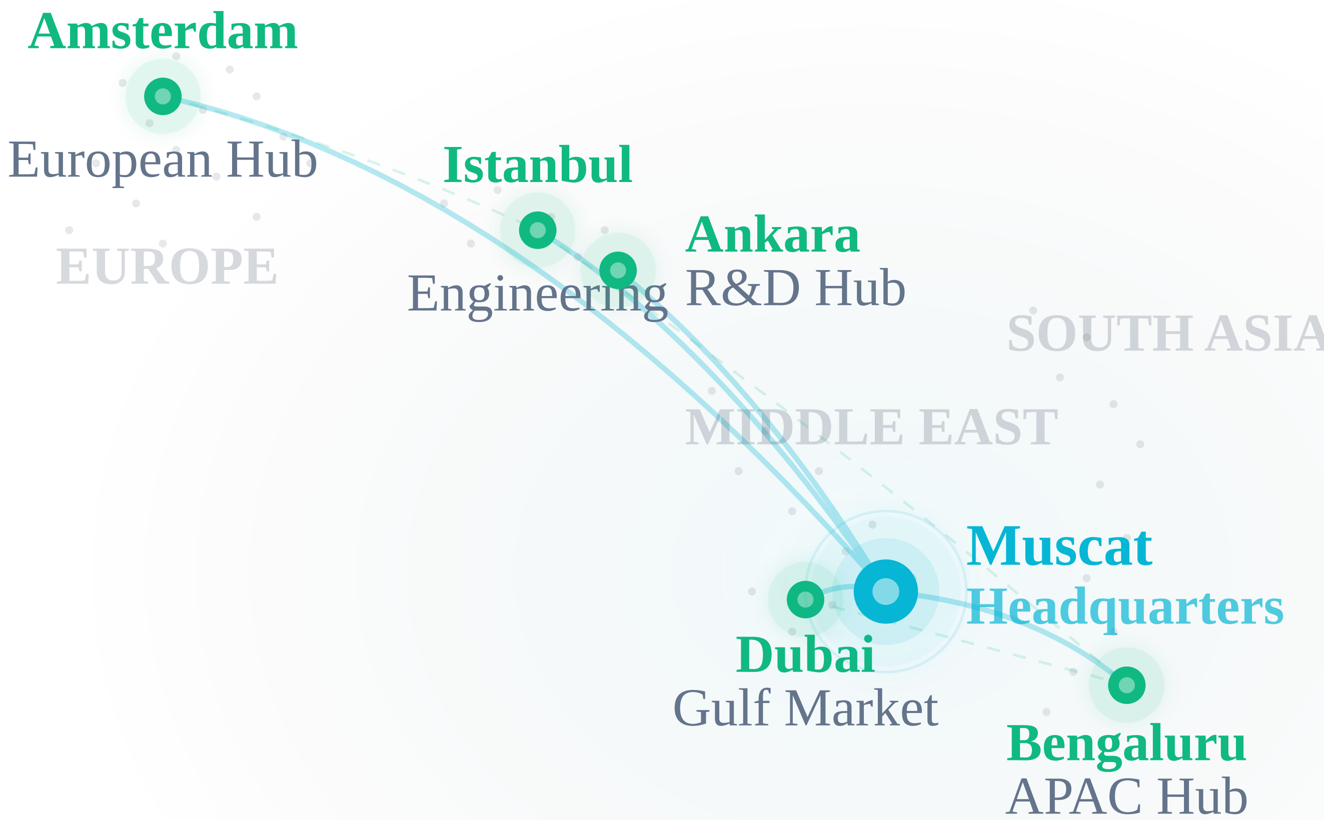

6 Cities | **3 Continents** | **24/7 Coverage**

Six cities. Three continents. Follow-the-sun 24/7 coverage. Deep expertise in banking, regulated industries, and Kubernetes.

Financial Impact

ROI & Total Cost of Ownership

Replace fragmented vendor contracts with a single platform subscription. Infrastructure savings start from day one.

60-85%

Projected infrastructure savings

65-75%

Projected staffing reduction

< 3 mo

Estimated payback period

Legacy IT Cost Structure vs OpenOva

Fragmented Legacy Stack

Cloud Infra
IaaS / PaaS
30%

Software
Licensing
25%

Observability
15%

Security
12%

Database
10%

Identity & Access — 5%

Support — 3%

~€1.8M / year (tools + infra)

vs

One Flat Platform
Fee

~€360K / yr

60-85% Total Cost Reduction
One flat platform fee replaces 7+ proprietary vendor contracts

Reference Case — ME Bank (~1,500 employees, 200+ cores, 2 regions)

Category	Before	With OpenOva	Saved
Observability	€250K	Included	100%
Virtualization	€200K	Included	100%
Database Licensing	€200K	Included	100%
Identity Platform	€60K	Included	100%
Platform Team (5 eng)	€600K	€240K (2 eng)	60%
SOC/NOC (5 eng)	€500K	€120K (1 + Specter)	76%
Total Annual	€1,810K	€360K+	80%

Why One Flat Fee Works

Legacy IT charges per-host, per-core, per-user, per-API-call — costs that scale linearly with growth. OpenOva replaces all of it with a single subscription: platform, observability, data, security, and identity included. Growth doesn't increase your platform bill.

80% average cost reduction. 7+ vendor contracts replaced by one flat fee. Payback in under three months.

Let's Build Your Platform

Pure open-source. AI-native. Zero lock-in.
From assessment to production in weeks, not years.

hello@openova.io · openova.io