





The Converged Cloud-Native Platform

Pure open-source. AI-native. Zero lock-in.

Infrastructure, data, AI, and communication — deployed in hours, not years.

Pure Open Source AI-Native CNCF Mature Community = Enterprise

openova.io · hello@openova.io · 2026

01 / 19






Why
Now

The \$6.15 Trillion Question

Gartner · Deloitte · MIT Sloan · CNBC
· Zylo · Flexera

\$6.15T Global IT Spending — 2026



The \$6.15 Trillion Breakdown

Gartner's February 2026 forecast projects global IT spending at **\$6.15 trillion** — a 9.8% increase year-over-year. The number represents the largest technology investment in human history, distributed across IT services (\$1.9T), enterprise software (\$1.4T), telecom services (\$1.4T), devices (\$836B), and data center infrastructure (\$653B). Yet beneath this record spending, the actual return

Why 75 Cents Buys Nothing New

The Deloitte CIO Survey reveals a devastating ratio: **75% of IT budgets** are consumed by operational maintenance. That's three-quarters of every dollar going to patching servers, renewing licenses, managing vendor contracts, servicing technical debt, and keeping legacy integrations from breaking. Only 25 cents funds anything new. Meanwhile, the average enterprise runs **660 SaaS**

The AI Readiness Crisis

AI isn't failing because the models are bad. **72% of CIOs** identify fragmented infrastructure as the #1 barrier to AI ROI (Gartner 2025). GPT-4, Claude, and Gemini perform brilliantly in isolation — but real enterprise deployment demands unified data, consistent APIs, and platform-wide observability. Instead, organizations have 660 disconnected tools, 3 competing cloud providers, and 14 different auth

Smart Money Exits Legacy

Wall Street's verdict is clear. **\$1 trillion** in software market cap was destroyed in Q1 2026 as investors concluded the old guard cannot deliver AI-ready infrastructure (CNBC). A single January trading day wiped **\$285 billion** from AI sector stocks (WinBuzzer). VMware's post-acquisition license restructuring hit customers with increases up to **+1,500%** (CIO/Avasant) — extraction, not innovation. PM SaaS **-58%**

on investment is collapsing. Enterprises are buying more tools, more platforms, and more infrastructure than ever before — but getting less value from each dollar.

applications, most adopted without centralized oversight. Zylo's 2025 SaaS Management Index found **53% of licenses are completely unused**. That's not a rounding error — it's systemic waste at industrial scale.

systems. The result: **95% of GenAI pilots** fail to deliver measurable business value (MIT Sloan 2025). Cloud waste adds insult: **27% of \$450B+** in global cloud spend is idle VMs and forgotten resources — ~\$17M/yr per org (Flexera).

YTD, IT outsourcing **-42%**, CRM **-38%**. Hedge funds made **\$24 billion** shorting legacy vendors in five weeks. The message is unambiguous: fragmented infrastructure is a liability.


Option B



The Industry Response

Five Megatrends. One Convergence.

Five megatrends — each backed by billions in investment and analyst consensus — all point to one architecture.



OpenOva sits at the intersection of all five. A full-stack, open-source, AI-native platform — with enterprise support. Born cloud-native. Zero lock-in. 03 / 19



Platform Architecture

Building Blocks — Concept View

Ten building blocks, one converged ecosystem. Each solves a distinct infrastructure challenge — independently adoptable, collectively transformative.

One platform, ten domains. Each building block solves a distinct infrastructure challenge — security, networking, observability, data, AI, communication, storage, scaling, GitOps, and lifecycle. Together, they replace the 660-app sprawl typical of large enterprises (10K+ employees — Zylo 2025) with a single converged ecosystem.

Fully pluggable — keep what works. OpenOva doesn't demand rip-and-replace. Already running a database you love? Keep it. Have a working CI/CD pipeline? Plug it in. Each building block is independently adoptable. Start with what hurts most, expand when ready. Your existing investments are preserved.

Containers first, VMs as a last resort. 95% of new digital workloads deploy on cloud-native platforms (Gartner 2025). OpenOva is container-native from the ground up. For the rare legacy workload that can't be containerized, Cloud CRDs (Crossplane) provision and manage VMs through the same GitOps workflow — no separate toolchain needed.

Born cloud-native, AI-native. Every component runs as a Kubernetes operator. GitOps-driven. Observable from day one. AI-ready by default. This is the platformization megatrend — purpose-built, not bolted on.

Convergence without the integration tax. What starts as à la carte adoption converges into a unified ecosystem — shared identity, shared observability, shared security — zero glue code, zero custom integrations.

AI needs clean infrastructure. 72% of CIOs can't get ROI from AI — not because the models are bad, but because fragmented data, siloed observability, and disconnected security prevent AI from delivering. OpenOva unifies the substrate so AI can actually work.

Security
TLS Automation Secrets Sync Secrets Vault Vuln Scan Runtime Security Supply Chain Trust
SBOM & CVE Policy Engine IAM
AI Hub
Model ServingServerlessLLM Inference Vector DBGraph DBEmbeddings Chat UIAI
GuardrailsLLM Observability LLM Gateway
Data & Integration
PostgreSQLMongoDB WireRedis Cache Event StreamingAnalytics DBStream Processing
Workflow EngineCDCData Lakehouse BI & DashboardsOpen Banking
Communication
Email ServerVideo & AudioWebRTC GWChat ProtocolPush Notify
GitOps & IaC
GitOps EngineGit ServerIaCCloud CRDs
Networking & Mesh
CNI & MeshL7 ProxyWAFDNS Sync GSLBReverse TunnelMesh VPNIPsec Gateway
Scaling & Resilience
Vertical ScalingEvent ScalingConfig ReloadHA Orchestration
Storage & Registry
Object StorageBackup & RestoreContainer Registry
AIOps
AIOps Brain Dashboards Telemetry Agent Log Aggregation Metrics Store Distributed Tracing
Instrumentation Search & Analytics Chaos Engineering Usage Metering
Bootstrap & Lifecycle
Initial Setup Day-2 Operations
Kubernetes (K3s)
Huawei Cloud
AWS
Oracle OCI
Hetzner

Ten building blocks. One converged ecosystem. Start with what hurts most — expand when ready. 04 / 19



Building Blocks — Tech View

Every component is real upstream open source — CNCF-graduated, battle-tested, community-backed. No proprietary forks. No rebranded wrappers.

\$8.8 trillion in demand-side value. A 2024 Harvard Business School study calculated that open-source software creates \$8.8T in demand-side economic value globally — meaning if organizations had to rebuild or license equivalent proprietary software, the cost would be 3.5× higher. 85% of financial services firms are increasing OSS adoption (FINOS 2025). 20% already save >\$1M/year.

CNCF-graduated, enterprise-proven. Every OpenOva component — Cilium, Falco, Flux, Grafana, Strimzi, KServe, Keycloak — is selected for production maturity and backed by 15.6M cloud-native developers. No rebranded forks. No proprietary wrappers. Real upstream projects.

Enterprise support, zero lock-in. Open source doesn't mean unsupported. OpenOva provides 24/7 enterprise support, SLAs, and expert guidance — while you keep full ownership of every line of code. Switch providers without switching technology. Community edition = Enterprise edition.

Open source is more secure, not less. The largest breaches of the decade all came from proprietary, closed-source software: a supply-chain attack on a network monitoring platform (\$100B+ impact), an endpoint agent update that crashed 8.5M devices worldwide, a file-transfer zero-day exploited across thousands of organizations (\$10B+ damages), a password manager breach exposing encrypted vaults, and a collaboration platform exploit affecting government agencies. Meanwhile, open-source code is audited by millions, with CVEs patched in hours. Intentional backdoors can't hide in code the world can read.

GUARDIAN

cert-manager ESO OpenBao Trivy Falco Sigstore Syft+Grype Kyverno Keycloak

CORTEX

KServeKnativevLLM MilvusNeo4jBGE LibreChatNeMoLangFuse Axon

FABRIC

CNPGFerretDBValkey StrimziClickHouseFlink TemporalDebeziumIceberg SupersetFingate

RELAY

StalwartLiveKitSTUNnerMatrixNtfy

PILOT

FluxGiteaOpenTofuCrossplane

SPINE

CiliumEnvoyCorazaExternalDNS k8gbfrpcNetBirdstrongSwan

SURGE

VPAKEDAReloaderContinuum

SILO

MinIOVeleroHarbor

INSIGHTS

Specter Grafana Alloy Loki Mimir Tempo OpenTelemetry OpenSearch Litmus OpenMeter

CATALYST

Bootstrap Wizard Lifecycle Manager

Kubernetes (K3s)

Huawei Cloud

AWS

Oracle OCI

Hetzner

Community IS enterprise. Same code, same features, same security — with 24/7 commercial support. 05 / 19



Portfolio Spotlight

CATALYST: Bootstrap & Lifecycle

From initial platform setup to Day-2 operations — guided, automated, and continuously managed. Runs on K3s across any cloud.

Platform engineering is the #1 strategic technology trend (Gartner 2024), yet 94% of core modernization projects exceed timelines and budgets (IBM IBV). The bottleneck isn't technology — it's the operational complexity of assembling, configuring, and maintaining dozens of cloud-native components across environments.

CATALYST eliminates this complexity. A guided bootstrap wizard provisions a production-ready platform in hours, not months. Day-2 lifecycle management — automated upgrades, health monitoring, drift repair — keeps every component current and healthy without manual intervention. Deploy on any cloud or bare metal, from edge to enterprise.



Bootstrap Wizard

Guided platform provisioning with interactive configuration and automated deployment



Lifecycle Manager

Continuous Day-2 operations with automated upgrades, health checks, and drift repair



K3s

Lightweight Kubernetes distribution optimized for edge and production workloads



Cloud Providers

Multi-cloud support: Huawei Cloud, AWS, Oracle OCI, Hetzner, and bare metal

Production-grade Kubernetes deployed in hours, not months. Automated Day-2 operations from day one. 06 / 19



Portfolio Spotlight

PILOT: GitOps & IaC

Git-driven delivery and infrastructure as code. Every change is auditable, reversible, and automated.

The shift to GitOps is accelerating — 76% of organizations now use or evaluate GitOps practices (CNCF 2025). Yet many rely on proprietary Git hosting and IaC tools with restrictive licensing changes. License compliance risk is real: the BSL/SSPL trend forces enterprises to re-evaluate their entire delivery toolchain.

PILOT delivers a fully integrated, license-safe GitOps stack. Every infrastructure change is Git-versioned, auditable, and reversible. Crossplane Cloud CRDs extend GitOps beyond Kubernetes — managing cloud resources, VMs, and databases through the same declarative workflow. Zero vendor lock-in on your delivery pipeline.



Flux

Continuous GitOps delivery engine that reconciles cluster state with Git as the single source of truth



Gitea

Lightweight internal Git hosting with built-in CI/CD pipelines and issue tracking



OpenTofu

MPL 2.0 licensed infrastructure-as-code for declarative cloud provisioning and management



Crossplane

Kubernetes-native cloud resource management using custom resource definitions and reconciliation

Every change auditable, reversible, and Git-driven. Infrastructure as code is not optional — it's the foundation. 07 / 19





SPINE: Networking & Mesh

eBPF-powered service mesh, L7 proxy, WAF, DNS automation, global load balancing, and secure connectivity.

Enterprise networking is fragmenting across sidecar proxies, proprietary load balancers, managed WAFs, and per-seat VPN licenses. Each adds latency, cost, and operational burden. Traditional service meshes inject sidecar containers into every pod — adding 15-30% memory overhead and significant complexity.

SPINE leverages eBPF to eliminate sidecars entirely. Cilium provides CNI, service mesh, mTLS, and L7 observability in a single kernel-level data plane. Add WAF protection, automated DNS, global load balancing, and secure connectivity across hybrid and edge environments — all managed declaratively through Kubernetes.



Cilium
eBPF-powered CNI with built-in service mesh, mTLS, and L7 visibility



Envoy
High-performance L7 proxy for traffic management, load balancing, and observability



Coraza
Open-source web application firewall implementing OWASP Core Rule Set protection



ExternalDNS
Automated DNS record synchronization between Kubernetes resources and DNS providers



k8gb
Global server load balancing with authoritative DNS for multi-cluster traffic management



frpc
Reverse tunnel proxy for exposing services behind NAT in hybrid and edge deployments



NetBird
Zero-config WireGuard mesh VPN for secure peer-to-peer connectivity across clusters



strongSwan
IPsec VPN gateway for site-to-site encryption and legacy network integration
eBPF-native networking at kernel speed. No sidecars, no overhead — from CNI to service mesh to VPN in one stack. 08 / 19



GUARDIAN: Security & Identity

End-to-end security: TLS, secrets, scanning, runtime protection, supply chain integrity, policy enforcement, and identity management.

The largest breaches of the decade came from proprietary, closed-source software — a supply-chain attack on a network monitoring platform (\$100B+ impact), an endpoint agent update that crashed 8.5M devices, a file-transfer zero-day (\$10B+ damages). Meanwhile, open-source code is audited by millions. Backdoors can't hide in code the world can read.

GUARDIAN wraps your platform in defense-in-depth: automated TLS, zero-trust secrets, eBPF runtime detection, supply chain signing, SBOM tracking, policy-as-code, and enterprise IAM. Every layer is open source, auditable, and integrated — from build pipeline to production runtime to identity.



cert-manager

Automates TLS certificate lifecycle with Let's Encrypt and private CA integration



ESO

Syncs secrets from external vaults into Kubernetes with automatic rotation support



OpenBao

MPL 2.0 secrets vault with per-cluster isolation and dynamic credential generation



Trivy

Comprehensive vulnerability scanner for container images, filesystems, and IaC templates



Falco

eBPF-based runtime threat detection for containers and Kubernetes workloads



Sigstore

Keyless container signing and verification for supply chain integrity assurance



Syft+Grype

Automated SBOM generation paired with vulnerability matching and CVE tracking



Kyverno

Policy engine for validating, mutating, and generating Kubernetes resources at admission



Keycloak

Enterprise identity and access management with OIDC, SAML, and FAPI support

Defense-in-depth from build to runtime. Zero-trust by default — mTLS, RBAC, policy-as-code, fully auditable. 09 / 19



Portfolio Spotlight

INSIGHTS: AIOps & Observability

Full observability stack with AI-native operations. Specter's 6 autonomous agents replace legacy SOAR/SIEM — DevSecOps, DevOps, SRE, FinOps, Compliance, and AI Ops.

Enterprises spend \$300-500K/yr on proprietary observability — per-host pricing that scales linearly with infrastructure. As cloud-native environments grow to thousands of containers, cost becomes prohibitive. Meanwhile, alert fatigue and siloed tools mean MTTR stays high despite massive spending.

INSIGHTS unifies logs, metrics, traces, and SIEM into a single stack with zero per-host fees. Specter's 6 autonomous AI agents — DevSecOps, SRE, FinOps, Compliance, DevOps, and AI Ops — replace the traditional SOC/NOC model. They detect, diagnose, and remediate issues before your team even wakes up. Chaos engineering and usage metering round out a complete operational platform.



Specter
Six autonomous AI agents for DevSecOps, SRE, FinOps, and continuous compliance



Grafana
Unified dashboards for metrics, logs, traces, and alerts in a single pane of glass



Alloy
Next-generation OpenTelemetry collector and telemetry pipeline agent by Grafana Labs



Loki
Horizontally scalable log aggregation system optimized for Kubernetes metadata queries



Mimir
Long-term metrics storage with global query capabilities and multi-tenancy support



Tempo
Distributed tracing backend with native OpenTelemetry support and deep Grafana integration



OpenTelemetry
Vendor-neutral instrumentation framework for automatic trace and metric collection



OpenSearch
Search and analytics engine powering hot-tier SIEM and security event correlation



Litmus
Chaos engineering platform for proactive resilience testing of Kubernetes workloads



OpenMeter
Real-time usage metering and billing engine for consumption-based pricing models
Six AI agents. Zero per-host fees. Detect, diagnose, and fix autonomously — unified logs, metrics, traces, and SIEM. 10 / 19



Portfolio Spotlight

SURGE: Scaling & Resilience

Intelligent autoscaling, configuration management, and high-availability orchestration.

Over-provisioning wastes 27% of cloud spend (\$44.5B globally), while under-provisioning causes outages. Traditional scaling is reactive — you pay for peak capacity 24/7 or risk downtime. Manual capacity planning fails in dynamic, event-driven architectures.

SURGE combines vertical right-sizing (VPA), event-driven horizontal scaling (KEDA), and continuous availability orchestration into an intelligent autoscaling engine. Workloads scale from 10 to 10,000 pods based on actual demand — queue depth, HTTP traffic, custom metrics. Zero-downtime deployments with automatic failover and configuration-change detection.



VPA
Automatically right-sizes pod CPU and memory requests based on actual utilization patterns



KEDA
Event-driven horizontal autoscaling triggered by queue depth, HTTP traffic, or custom metrics

↻
Reloader
Watches ConfigMaps and Secrets for changes and triggers rolling restarts automatically

∞
Continuum
Continuous availability orchestration ensuring zero-downtime deployments and failovers
Scale from 10 to 10,000 pods on demand. Smart autoscaling eliminates the \$44.5B global cloud waste problem. 11 / 19



Portfolio Spotlight

SILO: Storage & Registry

S3-compatible object storage, backup & disaster recovery, and container artifact registry.

Data sovereignty demands self-hosted storage. Cloud egress fees (\$0.08-0.12/GB) and proprietary registry lock-in make multi-cloud and hybrid deployments expensive and inflexible. Backup strategies that depend on a single cloud provider fail the sovereignty test.

SILO provides S3-compatible object storage, Kubernetes-native backup with cross-region replication, and a full container/Helm registry — all self-hosted. Run it on-prem, in your sovereign cloud, or alongside public clouds. Your data stays where regulations and business logic dictate. No egress surprises.



MinIO
S3-compatible high-performance object storage for cloud-native data and backup targets



Velero
Kubernetes-native backup, restore, and disaster recovery with cross-region replication



Harbor
Enterprise container and Helm chart registry with vulnerability scanning and RBAC
Your data stays where regulations demand. S3-compatible object storage, cross-region DR — zero egress surprises. 12 / 19



Portfolio Spotlight

CORTEX: AI Hub

Self-hosted AI infrastructure. Full control over models, prompts, and costs. No data leaves your environment.

\$2.5T will be spent on AI in 2026, with 90% going to inference (Gartner). Yet 95% of GenAI pilots fail ROI (MIT) — not because of the models, but because enterprises lack the infrastructure to run them. API-based AI services leak data, impose per-token fees, and create hard vendor lock-in on the most strategic technology of the decade.

CORTEX gives you full control. Self-hosted multi-model inference cuts token costs by 90%+. RAG pipelines with vector and graph databases. AI safety guardrails. LLM observability. A ChatGPT-like interface for business users. All running on your infrastructure, with your data, under your governance. Switch models freely — no API lock-in.



KServe
Kubernetes-native model serving with autoscaling, canary rollouts, and multi-framework support



Knative
Serverless platform for deploying event-driven and scale-to-zero inference workloads



vLLM
High-throughput inference engine with PagedAttention for optimal GPU utilization



Milvus
Distributed vector database for similarity search in RAG and recommendation pipelines



Neo4j
Graph database for knowledge graphs, entity relationships, and AI-powered reasoning



BGE
State-of-the-art embedding and reranking models for retrieval-augmented generation



LibreChat
ChatGPT-like interface supporting multiple LLM backends for enterprise AI chat



NeMo
AI safety guardrails framework for input/output filtering and responsible AI deployment



LangFuse
LLM observability platform with prompt tracing, cost analytics, and evaluation tools



Axon
Managed LLM gateway providing subscription-based inference access and rate limiting
Self-hosted AI inference cuts costs 90%+. Run any model on your infrastructure — no data leaves your environment. 13 / 19



Portfolio Spotlight

FABRIC: Data & Integration

Event-driven data platform: PostgreSQL, streaming, workflows, CDC, data lakehouse, and BI analytics.

Data infrastructure is the most expensive line item in enterprise IT. Proprietary database licensing alone costs \$100-300K/yr per organization, while managed event streaming adds another \$80-120K. Per-core pricing punishes scale. License audits punish growth. And each vendor creates its own data silo.

FABRIC replaces fragmented data infrastructure with a unified, open-source data platform. Cloud-native PostgreSQL with HA and automated backups. Event streaming for real-time architectures. Workflow orchestration for distributed transactions. A complete data lakehouse for analytics. Self-service BI. All integrated, all open source, all under your control.



CNPG
Cloud-native PostgreSQL operator with HA, automated backups, and declarative management



FerretDB
MongoDB wire protocol compatibility layer running on PostgreSQL for document workloads



Valkey
Redis-compatible in-memory cache and data store with full cluster mode support



Strimzi
Apache Kafka operator for production event streaming with automated topic management



ClickHouse
Columnar OLAP database for real-time analytics on billions of rows with SQL interface



Flink
Unified stream and batch processing engine for complex event processing pipelines



Temporal
Durable workflow orchestration for sagas, long-running processes, and distributed transactions



Debezium
Change data capture platform streaming database changes to Kafka in real time



Iceberg
Open table format enabling data lakehouse architecture with time-travel queries



Superset
Self-service BI platform with interactive dashboards, SQL Lab, and data exploration



Fingate
Open banking platform with FAPI authorization, OBIE APIs, and usage metering
PostgreSQL to real-time CDC to BI dashboards. Zero per-core licensing — one unified open-source data platform. 14 / 19



Portfolio Spotlight

FINGATE: Open Banking

End-to-end open banking solution compliant with FAPI and OBIE standards. Built entirely on open-source components from the OpenOva stack.

Open banking mandates (PSD2, OBIE, Saudi Arabia SAMA) require FAPI-certified APIs — yet proprietary platforms charge per-API-call fees that scale exponentially with transaction volume. Banks face a choice: pay ever-growing vendor fees or build from scratch. Neither option is sustainable for the 89% growth in ME/Africa digital banking.

FINGATE delivers a complete open banking solution built entirely on OpenOva's open-source stack. FAPI-certified authorization, OBIE-compliant APIs, real-time metering, payment orchestration with saga patterns, and financial-grade API protection. Zero per-API-call fees. Full regulatory compliance. Built for the next decade of digital finance.



Keycloak (FAPI)
Financial-grade API authorization server with Strong Customer Authentication for PSD2



OBIE APIs
Open Banking Implementation Entity standards for accounts, payments, and consent management



Temporal
Payment orchestration engine for PSD2 consent flows and settlement workflows

OpenMeter
API consumption metering with real-time billing and partner settlement reporting

Envoy + Coraza
API gateway with WAF protection, rate limiting, and financial-grade traffic management

CNPG
PostgreSQL for transactional ledger data with point-in-time recovery and replication

Strimzi + Debezium
Kafka event streaming with CDC for real-time balance synchronization and audit trails
FAPI-certified. PSD2 and SAMA compliant. 100% open source, zero per-API-call fees — built for the ME/Africa boom. 15 / 19



How We Work

Engagement Model

Two service tracks — platform and consultancy — take you from discovery to autonomous operations.



Platform Services
Subscription-based, self-service + managed

Community Edition

Full platform, self-managed. Community = Enterprise — same code, same features.

Enterprise Subscription Agreement (ESA)

24/7 enterprise support, SLAs, Expert Network access (80hr/quarter), continuous updates, quarterly business reviews.

Fully Managed SOC/NOC

Specter AI agents + human oversight. 24/7 autonomous operations — DevSecOps, SRE, FinOps, Compliance.

Pay-As-You-Go (Axon)

Managed LLM inference — consumption-based pricing, no GPU procurement, instant access.



Consultancy Services

Engagement-based, expert-led

Discovery & Strategy

Full landscape audit — we map every vendor contract, license obligation, and integration dependency. TCO/ROI financial modeling. Target architecture design that preserves what works. Board-ready strategy with clear milestones and risk quantification.

Statement of Work (SoW)

Fixed-scope delivery with defined milestones: platform deployment, workload containerization and migration (Exodus methodology), compliance alignment, hands-on team enablement, and knowledge transfer.

Time & Materials (T&M)

Embedded experts that scale with your transformation. Architecture reviews, performance optimization, security hardening, and specialist deep-dives — ramping up or down as phases progress.

CIO & DT Advisory

C-suite strategic partnership: technology roadmap aligned to business KPIs, organizational readiness assessment, change management, transition planning, and quarterly board reporting.

Every major legacy platform has a clear migration path. Production in hours. Majority cost eliminated from day one. 16 / 19



Global Reach

The Guild — Expert Network

Global presence with deep expertise in banking, regulated industries, and enterprise Kubernetes.

Muscat — Headquarters

Corporate headquarters, strategic operations, Middle East market

Dubai — Regional Presence

Financial services, enterprise partnerships, Gulf market

Amsterdam — European Hub

European market, regulatory compliance, enterprise sales

Istanbul — Engineering & Operations

Core development, platform engineering, DevOps, security

Ankara — Engineering Hub

Platform R&D, AI/ML engineering, product development

Bengaluru — Asia-Pacific Hub

Engineering scale, AI/ML talent, APAC market



Six cities. Three continents. Follow-the-sun 24/7 coverage. Deep expertise in banking, regulated industries, and Kubernetes.



ROI & Total Cost of Ownership

Replace fragmented vendor contracts with a single platform subscription. Infrastructure savings start from day one.

60-85%

Projected infrastructure savings


65-75%

Projected staffing reduction

< 3 mo

Estimated payback period

Legacy IT Cost Structure vs OpenOva





60-85% Total Cost Reduction
One flat platform fee replaces 7+ proprietary vendor contracts

Reference Case — ME Bank (~1,500 employees, 200+ cores, 2 regions)

Category	Before	With OpenOva	Saved
Observability	€250K	Included	100%
Virtualization	€200K	Included	100%
Database Licensing	€200K	Included	100%
Identity Platform	€60K	Included	100%
Platform Team (5 eng)	€600K	€240K (2 eng)	60%
SOC/NOC (5 eng)	€500K	€120K (1 + Specter)	76%
Total Annual	€1,810K	€360K+	80%

Why One Flat Fee Works

Legacy IT charges per-host, per-core, per-user, per-API-call — costs that scale linearly with growth. OpenOva replaces all of it with a single subscription: platform, observability, data, security, and identity included. Growth doesn't increase your platform bill.

80% average cost reduction. 7+ vendor contracts replaced by one flat fee. Payback in under three months. 18 / 19



Let's Build Your Platform

Pure open-source. AI-native. Zero lock-in.
From assessment to production in weeks, not years.

hello@openova.io · openova.io
19 / 19