






# The Converged Cloud-Native Platform

Pure open-source. AI-native. Zero lock-in.  
Infrastructure, data, AI, and communication — deployed in hours, not years.

Pure Open Source   AI-Native   CNCF Mature   Community = Enterprise  
openova.io · hello@openova.io · 2026

01 / 19





Why  
Now

## The \$6.15 Trillion Question

Gartner · Deloitte · MIT Sloan · CNBC ·  
Zylo · Flexera

### \$6.15T Global IT Spending — 2026



### The 75-Cent Problem

For every dollar an enterprise spends on IT, **75 cents** goes to keeping legacy systems alive: patching servers, renewing licenses, managing vendor contracts, servicing technical debt accumulated over decades. Only **25 cents** funds anything genuinely new. Your highest-paid engineers are

Gartner's February 2026 forecast projects global IT spending at **\$6.15 trillion**, a 9.8% increase over 2025. The number represents the largest technology investment in human history, spanning IT services (\$1.9T in consulting, outsourcing, and system integration), enterprise software (\$1.4T in SaaS and perpetual licenses), telecom services (\$1.4T), devices (\$836B), and data center systems (\$653B). Yet the return on each incremental dollar is collapsing.

plumbers,  
not  
architects.  
The Deloitte  
CIO Survey  
confirms  
this ratio  
has  
remained  
stubbornly  
unchanged  
for five  
consecutive  
years  
despite  
massive  
cloud  
migration  
efforts.

### The AI Readiness Crisis

AI isn't failing because the models are bad. **72% of CIOs** confirm fragmented infrastructure is the primary barrier to AI ROI (Gartner 2025). Enterprise deployment demands unified data, consistent APIs, and platform-wide observability. Instead, organizations have 660 disconnected tools, 3 cloud providers, and 14 auth systems. The devastating result: **95% of GenAI pilots** fail to deliver measurable business value (MIT Sloan 2025). Not bad algorithms — broken substrates.

### Enterprise Sprawl by the Numbers

The average large enterprise now operates **660 SaaS applications**, with orgs above 10,000 employees averaging \$284M in annual SaaS spend (Zylo 2025). Most adopted bottom-up with zero oversight: **53% of licenses completely unused**. Cloud waste adds **27% of \$450B+** in global spend — idle instances, over-provisioned VMs, forgotten resources totaling ~\$17M/yr per org burned (Flexera 2025). And **94% of modernization projects** exceed timelines and budgets (IBM IBV). Complexity kills transformation before it starts.

### Market Destruction — Q1 2026

Investors are punishing legacy vendors at historic levels. **\$1 trillion** in software market cap destroyed (CNBC). Single-day **\$285 billion** AI sector wipeout (WinBuzzer). VMware's post-acquisition license restructuring: **+1,500%** increases, forcing pay-or-migrate ultimatums on captive customers (CIO/Avasant). PM SaaS **-58%** YTD, IT outsourcing **-42%**, CRM **-38%**, virtualization **-27%**.


### Smart Money Speaks

Hedge funds made **\$24 billion** in profits shorting legacy enterprise vendors in just five weeks (CNBC). The message from capital markets is unambiguous: fragmented infrastructure is a liability, not an asset. Enterprises that continue bolting together dozens of incompatible tools face rising costs, falling productivity, and AI programs that never move beyond pilot stage. The question isn't whether to transform — it's how fast you can converge before your competitors do. The transformation window is closing.

|                        |                          |                          |
|------------------------|--------------------------|--------------------------|
| <b>-58%</b><br>PM SaaS | <b>-42%</b><br>IT Outsrc | <b>\$24B</b><br>Short \$ |
|------------------------|--------------------------|--------------------------|

## Five Megatrends. One Convergence.

Five megatrends — each backed by billions in investment and analyst consensus — all point to one architecture.



**OpenOva sits at the intersection of all five.** A full-stack, open-source, AI-native platform — with enterprise support. Born cloud-native. Zero lock-in. 03 / 19

## Building Blocks — Concept View

Ten building blocks, one converged ecosystem. Each solves a distinct infrastructure challenge — independently adoptable, collectively transformative.

**One platform, ten domains.** Each building block solves a distinct infrastructure challenge — security, networking, observability, data, AI, communication, storage, scaling, GitOps, and lifecycle. Together, they replace the 660-app sprawl typical of large enterprises (10K+ employees — Zylo 2025) with a single converged ecosystem.

**Fully pluggable — keep what works.** OpenOva doesn't demand rip-and-replace. Already running a database you love? Keep it. Have a working CI/CD pipeline? Plug it in. Each building block is independently adoptable. Start with what hurts most, expand when ready. Your existing investments are preserved.

**Containers first, VMs as a last resort.** 95% of new digital workloads deploy on cloud-native platforms (Gartner 2025). OpenOva is container-native from the ground up. For the rare legacy workload that can't be containerized, Cloud CRDs

(Crossplane) provision and manage VMs through the same GitOps workflow — no separate toolchain needed.

**Born cloud-native, AI-native.** Every component runs as a Kubernetes operator. GitOps-driven. Observable from day one. AI-ready by default. This is the platformization megatrend — purpose-built, not bolted on.

**Convergence without the integration tax.** What starts as à la carte adoption converges into a unified ecosystem — shared identity, shared observability, shared security — zero glue code, zero custom integrations.

**AI needs clean infrastructure.** 72% of CIOs can't get ROI from AI — not because the models are bad, but because fragmented data, siloed observability, and disconnected security prevent AI from delivering. OpenOva unifies the substrate so AI can actually work.

Security  
TLS Automation Secrets Sync Secrets Vault Vuln Scan Runtime Security Supply Chain Trust  
SBOM & CVE Policy Engine IAM

AI Hub  
Model ServingServerlessLLM Inference Vector DBGraph DBEmbeddings Chat UIAI  
GuardrailsLLM Observability LLM Gateway

Data & Integration  
PostgreSQLMongoDB WireRedis Cache Event StreamingAnalytics DBStream Processing  
Workflow EngineCDCData Lakehouse BI & DashboardsOpen Banking

Communication  
Email ServerVideo & AudioWebRTC GWChat ProtocolPush Notify

GitOps & IaC  
GitOps EngineGit ServerIaCCloud CRDs

Networking & Mesh  
CNI & MeshL7 ProxyWAFDNS Sync GSLBReverse TunnelMesh VPNIPsec Gateway

Scaling & Resilience  
Vertical ScalingEvent ScalingConfig ReloadHA Orchestration

Storage & Registry  
Object StorageBackup & RestoreContainer Registry

AIOps  
AIOps Brain Dashboards Telemetry Agent Log Aggregation Metrics Store Distributed Tracing  
Instrumentation Search & Analytics Chaos Engineering Usage Metering

Bootstrap & Lifecycle  
Initial Setup Day-2 Operations

Kubernetes (K3s)  
Huawei Cloud  
AWS  
Oracle OCI  
Hetzner

Ten building blocks. One converged ecosystem. Start with what hurts most — expand when ready. 04 / 19



Open Source Foundation

## Building Blocks — Tech View

Every component is real upstream open source — CNCF-graduated, battle-tested, community-backed. No proprietary forks. No rebranded wrappers.

**\$8.8 trillion in demand-side value.** A 2024 Harvard Business School study calculated that open-source software creates \$8.8T in demand-side economic value globally — meaning if organizations had to rebuild or license equivalent proprietary software, the cost would be 3.5× higher. 85% of financial services firms are increasing OSS adoption (FINOS 2025). 20% already save >\$1M/year.

**CNCF-graduated, enterprise-proven.** Every OpenOva component — Cilium, Falco, Flux, Grafana, Strimzi, KServe, Keycloak — is selected for production maturity and backed by 15.6M cloud-native developers. No rebranded forks. No

proprietary wrappers. Real upstream projects.

**Enterprise support, zero lock-in.** Open source doesn't mean unsupported. OpenOva provides 24/7 enterprise support, SLAs, and expert guidance — while you keep full ownership of every line of code. Switch providers without switching technology. Community edition = Enterprise edition.

**Open source is more secure, not less.** The largest breaches of the decade all came from proprietary, closed-source software: a supply-chain attack on a network monitoring platform (\$100B+ impact), an endpoint agent update that crashed 8.5M devices worldwide, a file-transfer zero-day exploited across thousands of organizations (\$10B+ damages), a password manager breach exposing encrypted vaults, and a collaboration platform exploit affecting government agencies. Meanwhile, open-source code is audited by millions, with CVEs patched in hours. Intentional backdoors can't hide in code the world can read.

GUARDIAN

cert-manager ESO OpenBao Trivy Falco Sigstore Syft+Grype Kyverno Keycloak

CORTEX

KServeKnativevLLM MilvusNeo4jBGE LibreChatNeMoLangFuse Axon

FABRIC

CNPGFerretDBValkey StrimziClickHouseFlink TemporalDebeziumIceberg SupersetFingate

RELAY

StalwartLiveKitSTUNnerMatrixNtfy

PILOT

FluxGiteaOpenTofuCrossplane

SPINE

CiliumEnvoyCorazaExternalDNS k8gbfrpcNetBirdstrongSwan

SURGE

VPAKEDAReloaderContinuum

SILO

MinIOVeleroHarbor

INSIGHTS

Specter Grafana Alloy Loki Mimir Tempo OpenTelemetry OpenSearch Litmus OpenMeter

CATALYST

Bootstrap Wizard Lifecycle Manager

Kubernetes (K3s)

Huawei Cloud

AWS

Oracle OCI

Hetzner

Community IS enterprise. Same code, same features, same security — with 24/7 commercial support. 05 / 19



Portfolio Spotlight

## CATALYST: Bootstrap & Lifecycle

From initial platform setup to Day-2 operations — guided, automated, and continuously managed. Runs on K3s across any cloud.

Platform engineering is the #1 strategic technology trend (Gartner 2024), yet 94% of core modernization projects exceed timelines and budgets (IBM IBV). The bottleneck isn't technology — it's the operational complexity of assembling, configuring, and maintaining dozens of cloud-native components across environments.

CATALYST eliminates this complexity. A guided bootstrap wizard provisions a production-ready platform in hours, not months. Day-2 lifecycle management — automated upgrades, health monitoring, drift repair — keeps every component current and healthy without manual intervention. Deploy on any cloud or bare metal, from edge to enterprise.



Bootstrap Wizard

Guided platform provisioning with interactive configuration and automated deployment



Lifecycle Manager

Continuous Day-2 operations with automated upgrades, health checks, and drift repair



K3s

Lightweight Kubernetes distribution optimized for edge and production workloads



Cloud Providers

Multi-cloud support: Huawei Cloud, AWS, Oracle OCI, Hetzner, and bare metal

Production-grade Kubernetes deployed in hours, not months. Automated Day-2 operations from day one. 06 / 19



Portfolio Spotlight

## PILOT: GitOps & IaC

Git-driven delivery and infrastructure as code. Every change is auditable, reversible, and automated.

The shift to GitOps is accelerating — 76% of organizations now use or evaluate GitOps practices (CNCF 2025). Yet many rely on proprietary Git hosting and IaC tools with restrictive licensing changes. License compliance risk is real: the BSL/SSPL trend forces enterprises to re-evaluate their entire delivery toolchain.

PILOT delivers a fully integrated, license-safe GitOps stack. Every infrastructure change is Git-versioned, auditable, and reversible. Crossplane Cloud CRDs extend GitOps beyond Kubernetes — managing cloud resources, VMs, and databases through the same declarative workflow. Zero vendor lock-in on your delivery pipeline.



Flux

Continuous GitOps delivery engine that reconciles cluster state with Git as the single source of truth



Gitea

Lightweight internal Git hosting with built-in CI/CD pipelines and issue tracking



OpenTofu

MPL 2.0 licensed infrastructure-as-code for declarative cloud provisioning and management



Crossplane

Kubernetes-native cloud resource management using custom resource definitions and reconciliation

Every change auditable, reversible, and Git-driven. Infrastructure as code is not optional — it's the foundation. 07 / 19



Portfolio Spotlight

## SPINE: Networking & Mesh

eBPF-powered service mesh, L7 proxy, WAF, DNS automation, global load balancing, and secure connectivity.

Enterprise networking is fragmenting across sidecar proxies, proprietary load balancers, managed WAFs, and per-seat VPN licenses. Each adds latency, cost, and operational burden. Traditional service meshes inject sidecar containers into every pod — adding 15-30% memory overhead and significant complexity.



SPINE leverages eBPF to eliminate sidecars entirely. Cilium provides CNI, service mesh, mTLS, and L7 observability in a single kernel-level data plane. Add WAF protection, automated DNS, global load balancing, and secure connectivity across hybrid and edge environments — all managed declaratively through Kubernetes.



Cilium  
eBPF-powered CNI with built-in service mesh, mTLS, and L7 visibility



Envoy  
High-performance L7 proxy for traffic management, load balancing, and observability



Coraza  
Open-source web application firewall implementing OWASP Core Rule Set protection



ExternalDNS  
Automated DNS record synchronization between Kubernetes resources and DNS providers



k8gb  
Global server load balancing with authoritative DNS for multi-cluster traffic management



frpc  
Reverse tunnel proxy for exposing services behind NAT in hybrid and edge deployments



NetBird  
Zero-config WireGuard mesh VPN for secure peer-to-peer connectivity across clusters



strongSwan  
IPsec VPN gateway for site-to-site encryption and legacy network integration  
eBPF-native networking at kernel speed. No sidecars, no overhead — from CNI to service mesh to VPN in one stack. 08 / 19



Portfolio Spotlight

## GUARDIAN: Security & Identity

End-to-end security: TLS, secrets, scanning, runtime protection, supply chain integrity, policy enforcement, and identity management.

The largest breaches of the decade came from proprietary, closed-source software — a supply-chain attack on a network monitoring platform (\$100B+ impact), an endpoint agent update that crashed 8.5M devices, a file-transfer zero-day (\$10B+ damages). Meanwhile, open-source code is audited by millions. Backdoors can't hide in code the world can read.

GUARDIAN wraps your platform in defense-in-depth: automated TLS, zero-trust secrets, eBPF runtime detection, supply chain signing, SBOM tracking, policy-as-code, and enterprise IAM. Every layer is open source, auditable, and integrated — from build pipeline to production runtime to identity.



cert-manager  
Automates TLS certificate lifecycle with Let's Encrypt and private CA integration



ESO  
Syncs secrets from external vaults into Kubernetes with automatic rotation support



OpenBao



 MPL 2.0 secrets vault with per-cluster isolation and dynamic credential generation

 Trivy  
Comprehensive vulnerability scanner for container images, filesystems, and IaC templates

 Falco  
eBPF-based runtime threat detection for containers and Kubernetes workloads

 Sigstore  
Keyless container signing and verification for supply chain integrity assurance

 Syft+Grype  
Automated SBOM generation paired with vulnerability matching and CVE tracking

 Kyverno  
Policy engine for validating, mutating, and generating Kubernetes resources at admission

 Keycloak  
Enterprise identity and access management with OIDC, SAML, and FAPI support  
Defense-in-depth from build to runtime. Zero-trust by default — mTLS, RBAC, policy-as-code, fully auditable. 09 / 19

 OpenOva



Portfolio Spotlight

## INSIGHTS: AIOps & Observability

Full observability stack with AI-native operations. Specter's 6 autonomous agents replace legacy SOAR/SIEM — DevSecOps, DevOps, SRE, FinOps, Compliance, and AI Ops.

Enterprises spend \$300-500K/yr on proprietary observability — per-host pricing that scales linearly with infrastructure. As cloud-native environments grow to thousands of containers, cost becomes prohibitive. Meanwhile, alert fatigue and siloed tools mean MTTR stays high despite massive spending.

INSIGHTS unifies logs, metrics, traces, and SIEM into a single stack with zero per-host fees. Specter's 6 autonomous AI agents — DevSecOps, SRE, FinOps, Compliance, DevOps, and AI Ops — replace the traditional SOC/NOC model. They detect, diagnose, and remediate issues before your team even wakes up. Chaos engineering and usage metering round out a complete operational platform.

 Specter  
Six autonomous AI agents for DevSecOps, SRE, FinOps, and continuous compliance

 Grafana  
Unified dashboards for metrics, logs, traces, and alerts in a single pane of glass

 Alloy  
Next-generation OpenTelemetry collector and telemetry pipeline agent by Grafana Labs

 Loki  
Horizontally scalable log aggregation system optimized for Kubernetes metadata queries

 Mimir  
Long-term metrics storage with global query capabilities and multi-tenancy support



Tempo  
Distributed tracing backend with native OpenTelemetry support and deep Grafana integration



OpenTelemetry  
Vendor-neutral instrumentation framework for automatic trace and metric collection



OpenSearch  
Search and analytics engine powering hot-tier SIEM and security event correlation



Litmus  
Chaos engineering platform for proactive resilience testing of Kubernetes workloads



OpenMeter  
Real-time usage metering and billing engine for consumption-based pricing models  
Six AI agents. Zero per-host fees. Detect, diagnose, and fix autonomously — unified logs, metrics, traces, and SIEM. 10 / 19



Portfolio Spotlight

## SURGE: Scaling & Resilience

Intelligent autoscaling, configuration management, and high-availability orchestration.

Over-provisioning wastes 27% of cloud spend (\$44.5B globally), while under-provisioning causes outages. Traditional scaling is reactive — you pay for peak capacity 24/7 or risk downtime. Manual capacity planning fails in dynamic, event-driven architectures.

SURGE combines vertical right-sizing (VPA), event-driven horizontal scaling (KEDA), and continuous availability orchestration into an intelligent autoscaling engine. Workloads scale from 10 to 10,000 pods based on actual demand — queue depth, HTTP traffic, custom metrics. Zero-downtime deployments with automatic failover and configuration-change detection.



VPA  
Automatically right-sizes pod CPU and memory requests based on actual utilization patterns



KEDA  
Event-driven horizontal autoscaling triggered by queue depth, HTTP traffic, or custom metrics



Reloader  
Watches ConfigMaps and Secrets for changes and triggers rolling restarts automatically



Continuum  
Continuous availability orchestration ensuring zero-downtime deployments and failovers  
Scale from 10 to 10,000 pods on demand. Smart autoscaling eliminates the \$44.5B global cloud waste problem. 11 / 19



Portfolio Spotlight

## SILO: Storage & Registry

S3-compatible object storage, backup & disaster recovery, and container artifact registry.

Data sovereignty demands self-hosted storage. Cloud egress fees (\$0.08-0.12/GB) and proprietary registry lock-in make multi-cloud and hybrid deployments expensive and inflexible. Backup strategies that depend on a single cloud provider fail the sovereignty test.

SILO provides S3-compatible object storage, Kubernetes-native backup with cross-region replication, and a full container/Helm registry — all self-hosted. Run it on-prem, in your sovereign cloud, or alongside public clouds. Your data stays where regulations and business logic dictate. No egress surprises.



MinIO  
S3-compatible high-performance object storage for cloud-native data and backup targets



Velero  
Kubernetes-native backup, restore, and disaster recovery with cross-region replication



Harbor  
Enterprise container and Helm chart registry with vulnerability scanning and RBAC  
Your data stays where regulations demand. S3-compatible object storage, cross-region DR — zero egress surprises. 12 / 19



Portfolio Spotlight

## CORTEX: AI Hub

Self-hosted AI infrastructure. Full control over models, prompts, and costs. No data leaves your environment.

\$2.5T will be spent on AI in 2026, with 90% going to inference (Gartner). Yet 95% of GenAI pilots fail ROI (MIT) — not because of the models, but because enterprises lack the infrastructure to run them. API-based AI services leak data, impose per-token fees, and create hard vendor lock-in on the most strategic technology of the decade.

CORTEX gives you full control. Self-hosted multi-model inference cuts token costs by 90%+. RAG pipelines with vector and graph databases. AI safety guardrails. LLM observability. A ChatGPT-like interface for business users. All running on your infrastructure, with your data, under your governance. Switch models freely — no API lock-in.



KServe  
Kubernetes-native model serving with autoscaling, canary rollouts, and multi-framework support



Knative  
Serverless platform for deploying event-driven and scale-to-zero inference workloads



vLLM  
High-throughput inference engine with PagedAttention for optimal GPU utilization



Milvus  
Distributed vector database for similarity search in RAG and recommendation pipelines



Neo4j  
Graph database for knowledge graphs, entity relationships, and AI-powered reasoning



BGE  
State-of-the-art embedding and reranking models for retrieval-augmented generation



LibreChat  
ChatGPT-like interface supporting multiple LLM backends for enterprise AI chat



NeMo  
AI safety guardrails framework for input/output filtering and responsible AI deployment



LangFuse  
LLM observability platform with prompt tracing, cost analytics, and evaluation tools



Axon  
Managed LLM gateway providing subscription-based inference access and rate limiting  
Self-hosted AI inference cuts costs 90%+. Run any model on your infrastructure — no data leaves your environment. 13 / 19



Portfolio Spotlight

FABRIC: Data & Integration

Event-driven data platform: PostgreSQL, streaming, workflows, CDC, data lakehouse, and BI analytics.

Data infrastructure is the most expensive line item in enterprise IT. Proprietary database licensing alone costs \$100-300K/yr per organization, while managed event streaming adds another \$80-120K. Per-core pricing punishes scale. License audits punish growth. And each vendor creates its own data silo.

FABRIC replaces fragmented data infrastructure with a unified, open-source data platform. Cloud-native PostgreSQL with HA and automated backups. Event streaming for real-time architectures. Workflow orchestration for distributed transactions. A complete data lakehouse for analytics. Self-service BI. All integrated, all open source, all under your control.



CNPG  
Cloud-native PostgreSQL operator with HA, automated backups, and declarative management



FerretDB  
MongoDB wire protocol compatibility layer running on PostgreSQL for document workloads



Valkey  
Redis-compatible in-memory cache and data store with full cluster mode support



Strimzi  
Apache Kafka operator for production event streaming with automated topic management



ClickHouse  
Columnar OLAP database for real-time analytics on billions of rows with SQL interface



Flink  
Unified stream and batch processing engine for complex event processing pipelines



Temporal  
Durable workflow orchestration for sagas, long-running processes, and distributed transactions



Debezium  
Change data capture platform streaming database changes to Kafka in real time



Iceberg  
Open table format enabling data lakehouse architecture with time-travel queries



Superset  
Self-service BI platform with interactive dashboards, SQL Lab, and data exploration



Fingate  
Open banking platform with FAPI authorization, OBIE APIs, and usage metering  
PostgreSQL to real-time CDC to BI dashboards. Zero per-core licensing — one unified open-source data platform. 14 / 19



Portfolio Spotlight

FINGATE: Open Banking

End-to-end open banking solution compliant with FAPI and OBIE standards. Built entirely on open-source components from the OpenOva stack.

Open banking mandates (PSD2, OBIE, Saudi Arabia SAMA) require FAPI-certified APIs — yet proprietary platforms charge per-API-call fees that scale exponentially with transaction volume. Banks face a choice: pay ever-growing vendor fees or build from scratch. Neither option is sustainable for the 89% growth in ME/Africa digital banking.

FINGATE delivers a complete open banking solution built entirely on OpenOva's open-source stack. FAPI-certified authorization, OBIE-compliant APIs, real-time metering, payment orchestration with saga patterns, and financial-grade API protection. Zero per-API-call fees. Full regulatory compliance. Built for the next decade of digital finance.



Keycloak (FAPI)  
Financial-grade API authorization server with Strong Customer Authentication for PSD2



OBIE APIs  
Open Banking Implementation Entity standards for accounts, payments, and consent management



Temporal  
Payment orchestration engine for PSD2 consent flows and settlement workflows



OpenMeter  
API consumption metering with real-time billing and partner settlement reporting



Envoy + Coraza  
API gateway with WAF protection, rate limiting, and financial-grade traffic management



CNPG  
PostgreSQL for transactional ledger data with point-in-time recovery and replication



Strimzi + Debezium  
Kafka event streaming with CDC for real-time balance synchronization and audit trails  
FAPI-certified. PSD2 and SAMA compliant. 100% open source, zero per-API-call fees — built for the ME/Africa boom. 15 / 19



How We Work

Engagement Model

Two service tracks — platform and consultancy — take you from discovery to autonomous operations.



## Platform Services

Subscription-based, self-service + managed

### Community Edition

Full platform, self-managed. Community = Enterprise — same code, same features.

### Enterprise Subscription Agreement (ESA)

24/7 enterprise support, SLAs, Expert Network access (80hr/quarter), continuous updates, quarterly business reviews.

### Fully Managed SOC/NOC

Specter AI agents + human oversight. 24/7 autonomous operations — DevSecOps, SRE, FinOps, Compliance.

### Pay-As-You-Go (Axon)

Managed LLM inference — consumption-based pricing, no GPU procurement, instant access.



## Consultancy Services

Engagement-based, expert-led

### Discovery & Strategy

Full landscape audit — we map every vendor contract, license obligation, and integration dependency. TCO/ROI financial modeling. Target architecture design that preserves what works. Board-ready strategy with clear milestones and risk quantification.

### Statement of Work (SoW)

Fixed-scope delivery with defined milestones: platform deployment, workload containerization and migration (Exodus methodology), compliance alignment, hands-on team enablement, and knowledge transfer.

### Time & Materials (T&M)

Embedded experts that scale with your transformation. Architecture reviews, performance optimization, security hardening, and specialist deep-dives — ramping up or down as phases progress.

## CIO & DT Advisory

C-suite strategic partnership: technology roadmap aligned to business KPIs, organizational readiness assessment, change management, transition planning, and quarterly board reporting.

Every major legacy platform has a clear migration path. Production in hours. Majority cost eliminated from day one. 16 / 19



Global Reach

## The Guild — Expert Network

Global presence with deep expertise in banking, regulated industries, and enterprise Kubernetes.

### Muscat — [Headquarters](#)

Corporate headquarters, strategic operations, Middle East market

### Dubai — Regional Presence

Financial services, enterprise partnerships, Gulf market

### Amsterdam — European Hub

European market, regulatory compliance, enterprise sales

### Istanbul — Engineering & Operations

Core development, platform engineering, DevOps, security


### Ankara — Engineering Hub

Platform R&D, AI/ML engineering, product development

### Bengaluru — Asia-Pacific Hub

Engineering scale, AI/ML talent, APAC market





6 Cities

3 Continents

24/7 Coverage

Six cities. Three continents. Follow-the-sun 24/7 coverage. Deep expertise in banking, regulated industries, and Kubernetes. 17 / 19



Financial Impact

## ROI & Total Cost of Ownership

Replace fragmented vendor contracts with a single platform subscription. Infrastructure savings start from day one.

60-85%

Projected infrastructure savings

65-75%

Projected staffing reduction


< 3 mo

Estimated payback period

# Legacy IT Cost Structure vs OpenOva

OpenOva

## Fragmented Legacy Stack



~€1.8M / year (tools + infra)

vs  
>

One Flat Platform Fee

~€360K / yr



**60-85% Total Cost Reduction**  
One flat platform fee replaces 7+ proprietary vendor contracts

## Reference Case — ME Bank (~1,500 employees, 200+ cores, 2 regions)

| Category              | Before  | With OpenOva        | Saved |
|-----------------------|---------|---------------------|-------|
| Observability         | €250K   | Included            | 100%  |
| Virtualization        | €200K   | Included            | 100%  |
| Database Licensing    | €200K   | Included            | 100%  |
| Identity Platform     | €60K    | Included            | 100%  |
| Platform Team (5 eng) | €600K   | €240K (2 eng)       | 60%   |
| SOC/NOC (5 eng)       | €500K   | €120K (1 + Specter) | 76%   |
| Total Annual          | €1,810K | €360K+              | 80%   |

## Why One Flat Fee Works

Legacy IT charges per-host, per-core, per-user, per-API-call — costs that scale linearly with growth. OpenOva replaces all of it with a single subscription: platform, observability, data,

security, and identity included. Growth doesn't increase your platform bill.

80% average cost reduction. 7+ vendor contracts replaced by one flat fee. Payback in under three months. 18 / 19



## **Let's Build Your Platform**

Pure open-source. AI-native. Zero lock-in.  
From assessment to production in weeks, not years.

[hello@openova.io](mailto:hello@openova.io) · [openova.io](https://openova.io)

19 / 19